

Yu Xiu

(434) 422-2707 | xiuyutiffany@gmail.com | [LinkedIn](#) | [Github](#) | [Portfolio](#)

EDUCATION

San José State University

Master of Science in Computer Science; GPA: 3.6

Stanford University Summer Session

Visiting Graduate Student, Computer Science

San José State University

Bachelor of Science in Computer Science; GPA: 3.7

San José, CA

2022 – Dec 2024

Stanford, CA

Jun – Aug 2024

San José, CA

Dec 2020

EXPERIENCE

Darluwin AI Labs (Early-stage Startup)

Founder & Full-Stack AI Engineer

Bay Area, CA

Aug 2025 – Present

DaResume

Demo

- Built and launched a **full-stack resume-to-job matching product** with React, TypeScript, FastAPI, and Supabase, helping users identify **missing evidence** against job requirements.
- Built a **LangGraph-based scoring workflow** that breaks job descriptions into requirements, checks each one against the resume, and returns evidence-based verdicts. Parallelized up to 12 model calls with **asyncio** to reduce analysis latency.
- Built **semantic search and duplicate detection** over user projects and experience data using pgvector and OpenAI embeddings, routing matches above a 0.86 similarity threshold to **human review** to prevent incorrect merges and data loss.
- Optimized **LLM latency and cost** with shared-prefix prompt caching and hash-based result caching, and added 45-second model timeouts within a 90-second request deadline for more reliable analysis.
- Designed **on-demand AI workflows** for bullet rewriting and gap-evidence scanning, keeping them off the main analysis path and caching results after first use to reduce unnecessary model calls and API cost.

AI Tutoring Platform

Demo

- Designed and deployed an AI tutoring platform with **evaluation tracking**, accuracy reports, and **version-controlled experiments** using Git.
- Conducted **prompt engineering experiments** comparing five strategies across GPT-4o-mini and GPT-5.2, improving grading accuracy from 49% to 74%.
- Implemented **responsible AI practices** including content moderation via the OpenAI Moderation API and confidence-based fallback routing.
- Fine-tuned BERT and RoBERTa as cost-efficient alternatives to GPT, reducing inference cost by **95%** and latency by **10x** while maintaining classification accuracy.
- Built **data pipelines** for dataset cleaning, validation checks, and structured outputs for downstream ML workflows.

NASA Ames Research Center

Systems and Software Engineering Intern

Mountain View, CA

Oct 2023 – Apr 2024

- Processed and analyzed **Excel-based telemetry data** from automated fault injection tests, identifying system anomalies and idle components to optimize test pipeline efficiency by **20%**.
- Developed MATLAB/Simulink simulation models for propulsion systems, creating **clear documentation** for model validation and **traceability**.
- Collaborated with **scientists and engineers** to debug system issues, translating technical requirements from sequence diagrams into implementation.
- Created **data visualizations** (charts, tables) to communicate findings to **cross-functional stakeholders**, supporting engineering decisions.

TECHNICAL SKILLS

Languages: Python, TypeScript, JavaScript, SQL

AI / ML: LangGraph, LangChain, OpenAI API, RAG, embeddings, semantic search, model evaluation, BERT

Frameworks and Data: FastAPI, React, PostgreSQL, Supabase, pgvector, ChromaDB

Cloud and Tools: Vercel, Render, GitHub Actions, Git, Docker